

## 1. Iris Flower Classification

**Description:** Classify iris flowers into three species (setosa, versicolor, virginica) based on features like petal and sepal length/width using classification algorithms.

### Steps:

- Load the Iris dataset (available in sklearn.datasets).
  - Explore and visualize the data using Pandas and Matplotlib/Seaborn.
  - Split the dataset into training and testing sets.
  - Train a model using a classifier (e.g., Logistic Regression or KNN).
  - Evaluate model accuracy using confusion matrix and classification report.
- 

## 2. House Price Prediction (Linear Regression)

**Description:** Predict the price of a house based on features like area, number of bedrooms, location, etc.

### Steps:

- Collect or use a dataset (like the Boston housing dataset or custom CSV).
  - Clean and preprocess the data (handle missing values, encode categories).
  - Use exploratory data analysis to understand correlations.
  - Train a Linear Regression model.
  - Evaluate performance using metrics like Mean Squared Error (MSE) or  $R^2$  score.
- 

## 3. Titanic Survival Prediction

**Description:** Predict whether a passenger survived the Titanic disaster based on features like age, class, gender, etc.

### Steps:

- Load the Titanic dataset (available on Kaggle or Seaborn).
  - Perform data cleaning and feature engineering.
  - Convert categorical variables using One-Hot Encoding.
  - Train a Logistic Regression model.
  - Evaluate using accuracy, precision, recall, and confusion matrix.
- 

## 4. Handwritten Digit Recognition

**Description:** Classify digits (0–9) from images of handwritten digits using ML.

**Steps:**

- Load the digits dataset from `sklearn.datasets` or use MNIST.
  - Normalize image pixel values.
  - Train a classifier (like RandomForest or SVM).
  - Test with a separate test set.
  - Display predictions with actual digit images using `matplotlib`.
- 

## 5. Movie Recommendation System (Content-Based)

**Description:** Recommend movies to users based on movie descriptions or genres using cosine similarity.

**Steps:**

- Load a movie dataset with titles and descriptions (e.g., TMDb dataset).
- Preprocess text (cleaning, vectorization using TF-IDF).
- Compute cosine similarity matrix.
- Define a function to recommend similar movies based on input title.
- Display top 5 recommendations.